Popper 2.0: A Container-Native Workflow Execution Engine For Testing Complex Applications and Reproducing Scientific Explorations

Ivo Jimenez, Jayjeet Chakraborty, Arshul Mansoori, Quincy Wofford and Carlos Maltzahn

Abstract—Software containers allow users to "bring their own environment" to shared computing platforms, reducing the friction between system administrators and their users. In recent years, multiple container runtimes have arisen, each addressing distinct needs (e.g. Singularity, Podman, rkt, among others), and an ongoing effort from the Linux Foundation (Open Container Initiative) is standardizing the specification of Linux container runtimes. While containers solve a big part of the "dependency hell" problem, there are scenarios where multi-container workflows are not fully addressed by existing runtimes or workflow engines. Current alternatives require a full scheduler (e.g. Kubernetes), a scientific workflow engine (e.g. Pegasus), or are constrained in the type of logic that users can express (e.g. Docker-compose). Ideally, users should be able to express workflows with the same user-friendliness and portability of **Dockerfiles** (write once, run anywhere). In this article, we introduce "Popper 2.0" a container-native workflow execution engine that allows users to express complex workflows similarly to how they do it in other scientific workflow languages, but with the advantage of running in container runtimes, bringing portability and ease of use to HPC scenarios. Popper 2.0 cleanly separates the three main concerns that are common in HPC scenarios: experimentation logic, environment preparation, and system configuration. To exemplify the suitability of the tool, we present three different case studies where we take examples from Machine Learning and HPC experiments and turn them into Popper workflows.

```yaml
version: 1
steps:
- id: build
  uses: ./workflows/containerized/actions/normalmodes
  runs: ./workflows/containerized/scripts/build.sh
  env:
    PEVSL_MAKEFILE_IN: ./workflows/containerized/makeconf/pEVSL_mkl.in
    NORMALMODES_MAKEFILE_IN: ./workflows/containerized/makeconf/NormalModes_mkl.in
    NUM_BUILD_JOBS: 1

- id: test
  uses: ./workflows/containerized/actions/normalmodes
  runs: ./workflows/containerized/scripts/test.sh

- id: run
  uses: ./workflows/containerized/actions/normalmodes
  runs: ./workflows/containerized/scripts/run.sh
  env:
    MPI_NUM_PROCESSES: 1
    INPUT_DIR: submodules/NormalModes/demos/

- id: validate
  uses: popperized/bin/sh@master
  runs: ./workflows/containerized/scripts/validate.sh

- id: generate vtk
  uses: docker://popperized/octave:4.4
  args: ./workflows/containerized/scripts/visualCmain.m
```

Figure 1: An end-to-end example of a workflow. On the left we have the `.yml` file that defines the workflow. On the right, a pictorial representation of it.

I. INTRODUCTION

Scientists and researchers often leave experimental artifacts like scripts, datasets, configuration files, etc. These when accompanied by lack of proper documentation makes the reproduction of the experiment highly difficult. Even with proper documentation of the steps that one need to follow to reproduce the experiments and rebuild the outcomes, it becomes cubersome due to the differences in the environment in which the artifacts were developed and in which they are being reproduced.

For example, a researcher working in a Windows environment writes some Windows PowerShell scripts. Now, if some other researcher with access to a Linux machine only wants to run those scripts in correct order in their environment, it would become time taking and difficult. They will probably end up spending a huge amount of time setting up VM's, finding the appropriate OS and interpreting the correct order of execution of the steps, hence making the process highly ineffecient and time consuming, which it does not need to be.

Through this paper, we propose a methodology in which a complex experimental workflow is decomposed into several discrete steps and each step executes in a separate container. Doing this way, makes the workflow platform independent. Although Software (Linux) containers are a relatively old technology [1], it was not until recently, with the rise of Docker, that they entered mainstream territory [2]. Since then, this technology has transformed the way applications get deployed in shared infrastructures, with 25% of companies using this form of software deployment [3], and a market size projected to be close to 5B by 2023 [4]. Docker has been the *de facto* container runtime, with other container runtimes such as Singularity [5], Charliecloud [6] and Podman[1] having emerged. Since, these container runtimes are available for almost every well known operating systems and architectures, experiments can be reproduced easily using containerized workflows in almost any environment.

At this point, one might think "Why not use a single container

---

[1] https://github.com/containers/libpod

for the entire workflow?". There are scenarios where a single container image is not suitable for implementing workflows associated to complex application testing or validating scientific explorations [7]. For example, a workflow might involve executing two steps, both of them requiring conflicting versions of a language runtime (e.g. Python 2.7 **and** Python 3.6). In this scenario, users could resort to solving such conflicts with the use of package managers, but this defeats the purpose of containers, which is to *not* have to do this sort of thing inside a container. More generally, the more complex a container image definition gets (a more complex `Dockerfile`), the more "monolithic" it gets, and thus the less maintainable and reusable it is. On the other hand, if an experimentation pipeline can be broken down into finer granular units, we end up having pieces of logic that are easier to maintain and reuse.

The above problem has been addressed by several tools in the past in distinct ways. Some popular Kubernetes based workflow engines include Argo and Pachyderm which requires access to a fully provisioned Kubernetes cluster. Some stable and generic workflow engines include Taverna and Pegasus which use custom workflow defination languages and assume workflow engine deployment prior to executing workflows.

The main contribution of this paper is a tool called Popper, which follows a container-native strategy for building reproducible worflows. The tool is described in detail in section II. This paper also presents three case studies of how popper can be used to quickly reproduce complex workflows in different environments in section III. We also present a detailed comparison of Popper with existing workflow engines in section V.

II. POPPER 2.0

In the remaining of this section we briefly expand on choosing YAML as the worflow specification language, as well as the design and implementation of the workflow execution engine (Popper 2.0); in particular, we dive into the container engine and the resource manager API layer.

A. Architecture

Here we describe the architecute or the dfd of popper

B. YAML as the workflow defination language

Here we show how yml serves better than other config languages

C. Workflow execution engine

Here we describe components of the workflow execution engine

1) Command line interface (PopperCLI): Here we talk about the features that PopperCLI provides.

2) Workflow runner: Here we talk about the work of the Workflow Runner

3) Internal representation of the workflow: Here we talk about how is the workflow interpreted internally

4) Resource managers: Here we talk about the functions of the resource manager

5) Container runtimes: Here we talk about the container runtimes.

III. CASE STUDY

A. Background

1) Docker: Here we talk about docker

2) Singularity: here we talk about singularity

3) SLURM: Here we talk about slurm

4) Kubernetes: Here we talk about Kubernetes

5) CI Services: Here we talk about some CI services

B. Execution Scenarios

Here we describe the 3 different execution scenarios

1) Single-Node local workflow execution: Single node workflow execution in the local machine for development purposes

2) Workflow execution in the Cloud using Kubernetes: Escalating the workflow execution to a GPU enabled cluster in the cloud for production grade execution

3) Exascale workflow execution in SLURM clusters: Planetscale workflow execution is super computing environments

IV. RESULTS

A. System Resource Usage

B. Overheads

V. RELATED WORK

Here we talk about other workflow execution solutions and how popper differentiates to them.

1) Generic workflow execution engines:

2) Container native workflow execution engines:

3) Cloud native workflow execution engines:

4) Workflow defination languages:

VI. CONCLUSION

A. Benefits

B. Challenges

C. Learning Curve

VII. FUTURE WORK

REFERENCES

[1] P.B. Menage, "Adding generic process containers to the linux kernel," *Proceedings of the Linux symposium*, Citeseer, 2007, pp. 45–57.

[2] D. Bernstein, "Containers and cloud: From lxc to docker to kubernetes," *IEEE Cloud Computing*, vol. 1, 2014, pp. 81–84.

[3] Datadog, "8 surprising facts about real Docker adoption," *8 surprising facts about real Docker adoption*, Jun. 2018. Available at: https://www.datadoghq.com/docker-adoption/.

[4] MarketsAndMarkets, "Application Container Market worth 4.98 Billion USD by 2023," 2018. Available at: https://www.marketsandmarkets.com/PressReleases/application-container.asp.

[5] G.M. Kurtzer, V. Sochat, and M.W. Bauer, "Singularity: Scientific containers for mobility of compute," *PloS one*, vol. 12, 2017, p. e0177459.

[6] R. Priedhorsky and T. Randles, "Charliecloud: Unprivileged containers for user-defined software stacks in hpc," *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis*, ACM, 2017, p. 36.

[7] C. Zheng and D. Thain, "Integrating containers into workflows: A case study using makeflow, work queue, and docker," *Proceedings of the 8th International Workshop on Virtualization Technologies in Distributed Computing*, ACM, 2015, pp. 31–38.